

Multicore Programming Task 1

Used Google cloud service **Ubuntu 20.04** disturbance and using the SSH connection for this task.

Only for the memory access analyses, I used the **Valgrind** because Vtune does not give the permission for this analysis.

Used these link steps for adding the profile with **Intel VTune** :

<https://www.intel.com/content/www/us/en/docs/vtune-profiler/installation-guide/2024-2/package-managers.html>

QUESTION 1

Because of I am using Cloud service for this Project I have to write my remote computer .

With **lscpu**:

```
enginsena00@instance-20241018-060600:~$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
Address sizes:         46 bits physical, 48 bits virtual
CPU(s):                32
On-line CPU(s) list:   0-31
Thread(s) per core:    2
Core(s) per socket:    16
Socket(s):              1
NUMA node(s):          1
Vendor ID:              GenuineIntel
CPU family:             6
Model:                  106
Model name:             Intel(R) Xeon(R) CPU @ 2.60GHz
Stepping:               6
CPU MHz:                2600.018
BogoMIPS:               5200.03
Hypervisor vendor:     FVM
Virtualization type:    full
L1d cache:              768 KiB
L1i cache:              512 KiB
L2 cache:                20 MiB
L3 cache:                54 MiB
NUMA node0 CPU(s):     0-31
Vulnerability Gather data sampling: Not affected
Vulnerability Itlb multihit: Not affected
Vulnerability L1tf: Not affected
Vulnerability Mds: Not affected
Vulnerability Meltdown: Not affected
Vulnerability Mmio stale data: Not affected
Vulnerability Reg file data sampling: Not affected
Vulnerability Retbleed: Not affected
Vulnerability Spec rstack overflow: Not affected
Vulnerability Spec store bypass: Mitigation: Speculative Store Bypass disabled via prctl and seccomp
Vulnerability Spectre v1: Mitigation: usercopy/swapgs barriers and __user pointer sanitization
Vulnerability Spectre v2: Mitigation: Enhanced / Automatic IBRS; IBPB conditional; RSB filling; PRRSB-eIBRS SW sequence; BHI SW loop, KVM SW loop
```

It can be seen that my processor architecture is x86- based .

X86_64 architecture means that my processor can handle 64-bit instructions.

CPU number is 32 , represents the total number of virtual CPUs on my VM.

Cores per socket means that 16 physical core machine has got.

Threads per core meaning each core supports these number of threads.

Model name: Intel(R) Xeon(R) CPU @ 2.60GHz, is my server processor model.

L1,L2 and L3 caches are caches from smallest to the biggest storage but faster to the slowest caches.

QUESTION 2

a) Perform an analysis to see how much time the sequential version of sobel filter spends on various part of the code (e.g., in Intel Vtune that corresponds to “hotspots” analysis). Report the result of the analysis and explain which part of the program is the most time consuming

```
vtune -collect hotspots -result-dir vtune_hotspots_seq -- ./sobel_seq input1.jpg
```

➤ **CPU Time** ⓘ: **1.270s**

Total Thread Count: 1

Paused Time ⓘ: 0s

Function	Module	CPU Time ⓘ	% of CPU Time ⓘ
generic_convolve	sobel_seq	1.060s	83.5%
combine	sobel_seq	0.100s	7.9%
func@0x37100	libjpeg.so.8	0.052s	4.1%
jpeg_read_scanlines	libjpeg.so.8	0.030s	2.4%
func@0x42100	libjpeg.so.8	0.020s	1.6%
[Others]	libjpeg.so.8	0.008s	0.6%

**N/A is applied to non-summable metrics.*

It can be seen that generic_convolve function is the most time consuming at the serial program .

b) Perform an analysis to see how much memory “consumption” the sequential version of sobel filter has, and report the findings.

```
vtune -collect memory-consumption -result-dir vtune_memory_consumption_seq -- ./sobel_seq input1.jpg
```

Function	Memory Consumption	Allocation/Deallocation Delta	Allocations	Module
main	113.0 MB	113.0 MB	3	sobel_seq
load_jpeg	37.8 MB	37.7 MB	12	sobel_seq
store_jpeg	158.5 KB	0.0 B	15	sobel_seq
printf	1.0 KB	1.0 KB	1	sobel_seq
[Unknown]			1	[Unknown]

**N/A is applied to non-summable metrics.*

In general , main functions taken most part of the consumption.

c) Perform “memory access” analysis to identify cache misses and bandwidth usage of the sequential version of sobel filter, and report the findings.

```
valgrind --tool=cachegrind --cachegrind-out-file=cachegrind.out ./sobel_seq input1.jpg
```

```
cg_annotate cachegrind.out > cachegrind_report.txt
```

l1 cache:	32768 B, 64 B, 8-way associative
D1 cache:	49152 B, 64 B, 12-way associative
l1 cache:	58720256 B, 64 B, 14-way associative
Command:	./sobel_seq input1.jpg
Data file:	cachegrind.out
Events recorded:	Ir I1mr I1mc Dr D1mr D1mc Dv D1mw D1mw
Events shown:	Ir I1mr I1mc Dr D1mr D1mc Dv D1mw D1mw
Event sort order:	Ir I1mr I1mc Dr D1mr D1mc Dv D1mw D1mw
Thresholds:	0.1 100 100 100 100 100 100 100 100
Include dirs:	
User annotated:	
Auto-annotation:	off

Ir	I1mr	I1mc	Dr	D1mr	D1mc	Dv	D1mw	D1mw	
9,646,224,584	2,589	2,499	2,243,766,227	85,326,459	4,611,689	184,524,491	77,555,064	4,718,233	PROGRAM TOTALS

Ir	I1mr	I1mc	Dr	D1mr	D1mc	Dv	D1mw	D1mw	file: function
7,753,399,838	6	6	1,731,406,884	82,386,917	3,042,602	75,296,474	75,270,732	3,535,820	/home/enginsena00/sobel_seq.c:generic_convolve
1,017,385,395	17	17	301,440,993	1,177,829	1,111,316	37,690,433	590,200	588,734	/home/enginsena00/sobel_seq.c:main
819,842,444	697	671	193,072,798	1,694,116	454,782	52,704,352	1,467,122	591,426	???:???
29,953,670	5	5	4,081,292	25,200	49	2,489,006	369	0	???:jpeg_fill_bit_buffer
13,305,738	13	13	12,802,254	22,270	4	12,764,099	222,028	1,148	/build/glibc-LcI20x/glibc-2.31/string/./sysdeps/x86_64/multiarch/memmove-vec-unaligned-erms.S:memcpy_avx_unaligned_erms

I put the cache misses txt files into the reports folder.

QUESTION 3

Because of the generic_convolve function is the most time consuming ,I parallelized the loop part of this function.

My strategy is the dividing the sobel_X and sobel_Y edge threads as a two threads and apply this for multiple threads.

As a difference of the serial code,

I added the start and end rows for each thread.

Getting the thread number from command line:

```
if (argc >= 3) {
    num_threads = atoi(argv[2]);
    if (num_threads <= 0) {
        fprintf(stderr, "Invalid number of threads specified.\n");
        return -1;
    }
}
```

Pthread_create and join to wait all threads to finish.

```

    /* Create thread */
    pthread_create(&threads[i], NULL, thread_convolve, (void *)&thread_data[i]);
}

/* Wait for threads to finish */
for (i = 0; i < num_threads; i++) {
    pthread_join(threads[i], NULL);
}

```

Allocated the threads with malloc.

```

/* Allocating memory for threads and thread data structures */
pthread_t *threads = malloc(num_threads * sizeof(pthread_t));
struct thread_data_t *thread_data = malloc(num_threads * sizeof(struct thread_data_t));

```

It is necessary to calculate image dimensions for all the threads:

```

/* Calculating image dimensions */
int i;
int width = image.x * image.depth; /* Total number of bytes in a row */
int depth = image.depth; /* Number of color channels (e.g., 3 for
RGB) */
int total_rows = image.y - 2; /* Number of rows to process (excluding
borders) */
int rows_per_thread = total_rows / num_threads; /* Number of rows each thread will
process */
int remaining_rows = total_rows % num_threads; /* Extra rows to distribute among
threads */
int current_row = 1; /* Start from 1 to exclude border */

```

Then dividing the work among threads for sobel_x filter:

```

/* Divide the work among threads */
for (i = 0; i < num_threads; i++) {
    /* Initialize thread data */
    thread_data[i].old = &image;
    thread_data[i].new = &sobel_x;
    thread_data[i].filter = &sobel_x_filter;
    thread_data[i].depth = depth;
    thread_data[i].width = width;
    thread_data[i].start_row = current_row;
    thread_data[i].end_row = current_row + rows_per_thread;

    /* Distribute any remaining rows */
    if (remaining_rows > 0) {
        thread_data[i].end_row++;
        remaining_rows--;
    }
    current_row = thread_data[i].end_row;
}

```

And dividing the sobel_y filter work among threads too:

```

/* Dividing the work among threads */
for (i = 0; i < num_threads; i++) {
    /* Initialize thread data */
    thread_data[i].old = &image;
    thread_data[i].new = &sobel_y;
    thread_data[i].filter = &sobel_y_filter;
    thread_data[i].depth = depth;
    thread_data[i].width = width;
    thread_data[i].start_row = current_row;
    thread_data[i].end_row = current_row + rows_per_thread;

    /* Distribute any remaining rows */
    if (remaining_rows > 0) {
        thread_data[i].end_row++;
        remaining_rows--;
    }
    current_row = thread_data[i].end_row;
}

```

Lastly , freed the memory.

```

/* Free allocated memory */
free(image.pixels);
free(sobel_x.pixels);
free(sobel_y.pixels);
free(new_image.pixels);
free(threads);
free(thread_data);

```

QUESTION 4

- a) Perform an analysis to see how much time the parallel version of sobel filter spends on various functions (e.g., “hotspots” analysis in Vtune) when different number of threads are used. You should perform separate analysis for thread counts of 2, 4, 8, 16 and 32. Report the results similar to above.

For 2 Thread:

Function	Module	CPU Time ⓘ	% of CPU Time ⓘ
thread_convolve	sobel_par	0.600s	73.2%
combine	sobel_par	0.092s	11.2%
jpeg_read_scanlines	libjpeg.so.8	0.060s	7.3%
func@0x37100	libjpeg.so.8	0.032s	3.9%
jpeg_write_scanlines	libjpeg.so.8	0.020s	2.4%
[Others]	N/A*	0.016s	2.0%

*N/A is applied to non-summable metrics.

For 4 Thread:

Function	Module	CPU Time ?	% of CPU Time ?
thread_convolve	sobel_par	0.590s	72.8%
combine	sobel_par	0.102s	12.6%
jpeg_read_scanlines	libjpeg.so.8	0.038s	4.7%
func@0x37100	libjpeg.so.8	0.036s	4.4%
func@0x3e840	libjpeg.so.8	0.024s	3.0%
[Others]	N/A*	0.020s	2.5%

*N/A is applied to non-summable metrics.

For 8 Thread:

Function	Module	CPU Time ?	% of CPU Time ?
thread_convolve	sobel_par	0.530s	70.7%
combine	sobel_par	0.102s	13.6%
func@0x3e840	libjpeg.so.8	0.032s	4.3%
jpeg_read_scanlines	libjpeg.so.8	0.032s	4.3%
jpeg_write_scanlines	libjpeg.so.8	0.020s	2.7%
[Others]	N/A*	0.034s	4.5%

*N/A is applied to non-summable metrics.

For 16 Thread:

Function	Module	CPU Time ?	% of CPU Time ?
thread_convolve	sobel_par	0.530s	70.7%
combine	sobel_par	0.102s	13.6%
func@0x37100	libjpeg.so.8	0.036s	4.8%
func@0x42100	libjpeg.so.8	0.028s	3.7%
jpeg_write_scanlines	libjpeg.so.8	0.024s	3.2%
[Others]	N/A*	0.030s	4.0%

*N/A is applied to non-summable metrics.

For 32 Thread:

Function	Module	CPU Time ②	% of CPU Time ②
thread_convolve	sobel_par	0.700s	76.9%
combine	sobel_par	0.100s	11.0%
jpeg_read_scanlines	libjpeg.so.8	0.050s	5.5%
func@0x37100	libjpeg.so.8	0.040s	4.4%
jpeg_write_scanlines	libjpeg.so.8	0.020s	2.2%

**N/A is applied to non-summable metrics.*


We can see that increasing the thread counts, decreases the operation time.
This means that my parallelization is effectively works.

- b) Perform “threading” analysis for thread counts of 2, 4, 8, 16 and 32, to see how efficiently an application uses available cores. Report the results similar to above.

For 2 thread:

Effective CPU Utilization Histogram


This histogram displays a percentage of the wall time the specific number of CPUs were running simultaneously. Spin and Overhead time adds to the Idle CPU utilization value.



For 4 thread:

Effective CPU Utilization Histogram


This histogram displays a percentage of the wall time the specific number of CPUs were running simultaneously. Spin and Overhead time adds to the Idle CPU utilization value.



For 8 thread:

Effective CPU Utilization Histogram


This histogram displays a percentage of the wall time the specific number of CPUs were running simultaneously. Spin and Overhead time adds to the Idle CPU utilization value.



For 16 thread:

Effective CPU Utilization Histogram


This histogram displays a percentage of the wall time the specific number of CPUs were running simultaneously. Spin and Overhead time adds to the Idle CPU utilization value.



For 32 thread:

Effective CPU Utilization Histogram

This histogram displays a percentage of the wall time the specific number of CPUs were running simultaneously. Spin and Overhead time adds to the Idle CPU utilization value.



If we look at the CPU diagrams, CPU utilization peaks at the 2 threads using but with the 32 threads may be all of the threads are not actively running so not reach the peak and as we see it is decreasing.

- c) Perform “memory access” analysis for thread counts of 2, 4, 8, 16 and 32. Report the results similar to above.

For 2 thread:

```

I1 cache:      32768 B, 64 B, 8-way associative
D1 cache:      49152 B, 64 B, 12-way associative
LL cache:      58720256 B, 64 B, 14-way associative
Command:       ./sobel_par input1.jpg 2
Data file:     cachegrind_out_2
Events recorded: Ic I1mr ILmr Dr D1mr DLmr Dw D1mw DLmw
Events shown:   Ic I1mr ILmr Dr D1mr DLmr Dw D1mw DLmw
Event sort order: Ic I1mr ILmr Dr D1mr DLmr Dw D1mw DLmw
Thresholds:     0.1 100 100 100 100 100 100 100 100
Include dirs:
User annotated:
Auto-annotation: off

```

<u>Ic</u>	<u>I1mr</u>	<u>ILmr</u>	<u>Dr</u>	<u>D1mr</u>	<u>DLmr</u>	<u>Dw</u>	<u>D1mw</u>	<u>DLmw</u>
10,586,759,574	2,849	2,707	2,507,055,419	13,532,671	2,423,072	184,508,248	5,809,594	2,359,100
PROGRAM TOTALS								
<u>Ic</u>	<u>I1mr</u>	<u>ILmr</u>	<u>Dr</u>	<u>D1mr</u>	<u>DLmr</u>	<u>Dw</u>	<u>D1mw</u>	<u>DLmw</u>
<u>file: function</u>								
8,731,545,608	6	5	2,032,362,518	10,589,922	787,659	75,270,768	3,529,972	1,176,656
/home/enginsena00/sobel_par.c:thread_convolve								
979,706,917	24	24	263,756,517	1,177,473	1,177,469	37,690,481	590,201	588,733
/home/enginsena00/sobel_par.c:main								
819,842,712	734	698	193,072,925	1,698,058	454,800	52,704,380	1,462,159	
591,426	???:???							
29,953,670	5	5	4,081,292	25,200	49	2,489,006	689	
0	???:jpeg_fill_bit_buffer							
13,305,738	13	13	12,802,254	22,522	4	12,764,099	222,194	1,148

For 4 thread:

```

I1 cache:      32768 B, 64 B, 8-way associative
D1 cache:      49152 B, 64 B, 12-way associative
LL cache:      58720256 B, 64 B, 14-way associative
Command:       ./sobel_par input1.jpg 4
Data file:     cachegrind_out_4
Events recorded: Ic I1mr ILmr Dr D1mr DLmr Dw D1mw DLmw
Events shown:   Ic I1mr ILmr Dr D1mr DLmr Dw D1mw DLmw
Event sort order: Ic I1mr ILmr Dr D1mr DLmr Dw D1mw DLmw
Thresholds:     0.1 100 100 100 100 100 100 100 100
Include dirs:
User annotated:
Auto-annotation: off

```

<u>Ic</u>	<u>I1mr</u>	<u>ILmr</u>	<u>Dr</u>	<u>D1mr</u>	<u>DLmr</u>	<u>Dw</u>	<u>D1mw</u>	<u>DLmw</u>
10,586,764,836	2,849	2,706	2,507,056,821	13,522,556	2,189,690	184,509,224	5,808,412	2,359,200
PROGRAM TOTALS								
<u>Ic</u>	<u>I1mr</u>	<u>ILmr</u>	<u>Dr</u>	<u>D1mr</u>	<u>DLmr</u>	<u>Dw</u>	<u>D1mw</u>	<u>DLmw</u>
<u>file: function</u>								
8,731,545,808	6	4	2,032,362,586	10,589,961	588,092	75,270,804	3,529,984	1,176,656
/home/enginsena00/sobel_par.c:thread_convolve								
979,707,041	24	24	263,756,537	1,177,474	1,143,623	37,690,517	590,202	588,734
/home/enginsena00/sobel_par.c:main								
819,842,888	735	699	193,073,013	1,691,885	454,801	52,704,380	1,461,797	
591,425	???:???							
29,953,670	5	5	4,081,292	25,200	49	2,489,006	689	
0	???:jpeg_fill_bit_buffer							
13,305,738	13	13	12,802,254	18,554	4	12,764,099	221,281	1,149

For 8 thread:

```

I1 cache: 32768 B, 64 B, 8-way associative
D1 cache: 49152 B, 64 B, 12-way associative
LL cache: 58720256 B, 64 B, 14-way associative
Command: ./sobel_par input1.jpg 8
Data file: cachegrid_out_8
Events recorded: Ir I1mr I1Lmr Dr D1mr D1Lmr Dw D1mw D1Lmw
Events shown: Ir I1mr I1Lmr Dr D1mr D1Lmr Dw D1mw D1Lmw
Event sort order: Ir I1mr I1Lmr Dr D1mr D1Lmr Dw D1mw D1Lmw
Thresholds: 0.1 100 100 100 100 100 100 100
Include dirs:
User annotated:
Auto-annotation: off

```

Ir	I1mr	<u>I1mr</u>	Dr	D1mr	<u>D1mr</u>	<u>Dw</u>	D1mw	<u>D1mw</u>
10,586,781,019	2,863	2,720	2,507,061,286	13,523,845	2,138,487	184,511,951	5,801,884	2,359,
PROGRAM TOTALS								

Ir	I1mr	I1mr	Dr	D1mr	D1mr	Dw	D1mw	D1mw
8,731,546,288	6	4	2,032,362,722	10,590,038	561,264	75,270,876	3,530,006	1,176,656
/home/engine/na00/sobel	par.c:thread_convolve							
979,707,289	24	24	263,756,577	1,177,476	1,119,082	37,690,589	590,207	588,734
/home/engine/na00/sobel	par.c:main							
819,843,442	736	700	193,073,289	1,689,292	454,815	52,704,384	1,455,158	
591,425	???:???							
29,953,670	5	5	4,081,292	25,200	49	2,489,006	689	
0	???jpeg_fill_bit_buffer							
13,305,738	13	13	12,802,254	21,826	4	12,764,099	221,156	1,149

For 16 thread:

```

II cache:      32768 B, 64 B, 8-way associative
D1 cache:     49152 B, 64 B, 12-way associative
LL cache:     58720256 B, 64 B, 14-way associative
Command:       ./sobel_par input.jpg 16
Data file:     cachegrind_out_16
Events recorded: Ir I1mr I1mc Dr D1mr D1mc Dw D1mw D1mcw
Events shown:   Ir I1mr I1mc Dr D1mr D1mc Dw D1mw D1mcw
Event sort order: Ir I1mr I1mc Dr D1mr D1mc Dw D1mw D1mcw
Thresholds:     0.1 100 100 100 100 100 100 100 100
Include dirs:
User annotated:
Auto-annotation: off

```

<u>Ir</u>	<u>I1mr</u>	<u>I1mc</u>	<u>Dr</u>	<u>D1mr</u>	<u>D1mc</u>	<u>Dw</u>	<u>D1mw</u>	<u>D1mw</u>
10,586,810,758	2,865	2,721	2,507,069,406	13,545,453	2,046,857	184,516,909	5,811,704	2,359,561
PROGRAM TOTALS								

<u>Ir</u>	<u>I1mr</u>	<u>I1mr</u>	<u>Dr</u>	<u>D1mr</u>	<u>DLmr</u>	<u>Dw</u>	<u>D1mw</u>	<u>DLmw</u>
<u>file:function</u>								
8,731,547,008	6	4	2,032,362,994	10,590,215	505,767	75,271,020	3,530,056	1,176,658
/home/engineersna0	sobel	par.c:	thread_convolve					
979,707,785	24	24	263,756,657	1,177,480	1,082,824	37,690,733	590,216	588,734
/home/engineersna0	sobel	par.c:	main					
819,844,530	737	701	193,073,833	1,713,721	454,817	52,704,384	1,464,650	
591,422	???:???							
29,953,670	5	5	4,081,292	25,200	49	2,489,006	689	
0	???:jpeg_fill_bit_buffer							
13,305,738	13	13	12,802,254	18,581	4	12,764,099	220,961	1,148

For 32 thread:

```

11          32768 B, 64 B, 8-way associative
D1 cache: 49152 B, 64 B, 12-way associative
L1 cache: 5872056 B, 64 B, 14-way associative
Command: /sobel_par_input1.jpg 32
Data file: cacheind_out_32
Events recorded: Ic Ilm Ilmc Dr Dlmr Dllmc Dsc Dlmw Dllmw
Events shown: Ic Ilm Ilmc Dr Dlmr Dllmc Dsc Dlmw Dllmw
Event sort order: Ic Ilm Ilmc Dr Dlmr Dllmc Dsc Dlmw Dllmw
Thresholds: 0.1 100 100 100 100 100 100 100 100
Include dirs:
User annotated:
Auto-annotation: off

```

<u>Ir</u>	<u>I1mr</u>	<u>I1mc</u>	<u>Dr</u>	<u>D1mr</u>	<u>D1mc</u>	<u>Dw</u>	<u>D1mw</u>	<u>D1mw</u>	
10,586,871,665	2,867	2,721	2,507,085,771	13,531,046	2,183,985	184,526,888	5,819,220	2,359,946	PROGRAM TOTALS

[illegible]

Ir means instruction counts and actually it is increasing from 2 threads to the 32 threads.

I1 and IL means that level 1 and IL is the last level instruction caches.

D1 and DL is the data caches.

Mr-read misses

Mw-write misses

We can see that increasing the number of threads does not significantly much impact cache misses.